



## ZYRO HACKATHON

Team : PHEONIX

Mayaluri Anusha ([LinkedIn Profile](#))

Chintha Sai Ganesh ([LinkedIn Profile](#))

AI ML Problem Statement : IMAGE CLASSIFYING  
CHALLENGE

**Google Colab File Link :**

<https://colab.research.google.com/drive/14lrcuK1gSlx4dhD4wmLno9jUg7BsCYTr?usp=sharing>

**GitHub Repo with all code files :**

[https://github.com/MAYALURI-ANUSHA/Zyro\\_PHEONIX](https://github.com/MAYALURI-ANUSHA/Zyro_PHEONIX)

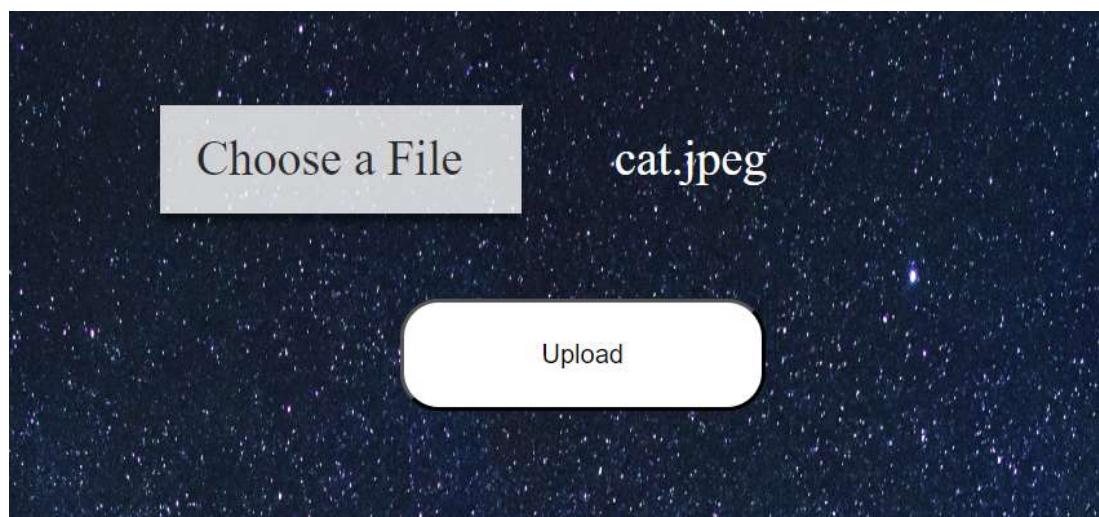
**Figma Design File :**

<https://www.figma.com/file/D3ZmWkJr0Kbd8mInveKXKK/Untitled?type=design&node-id=0-1&mode=design&t=jPnZa2BCjoHW90ai-0>

**Figma Prototype Link :**

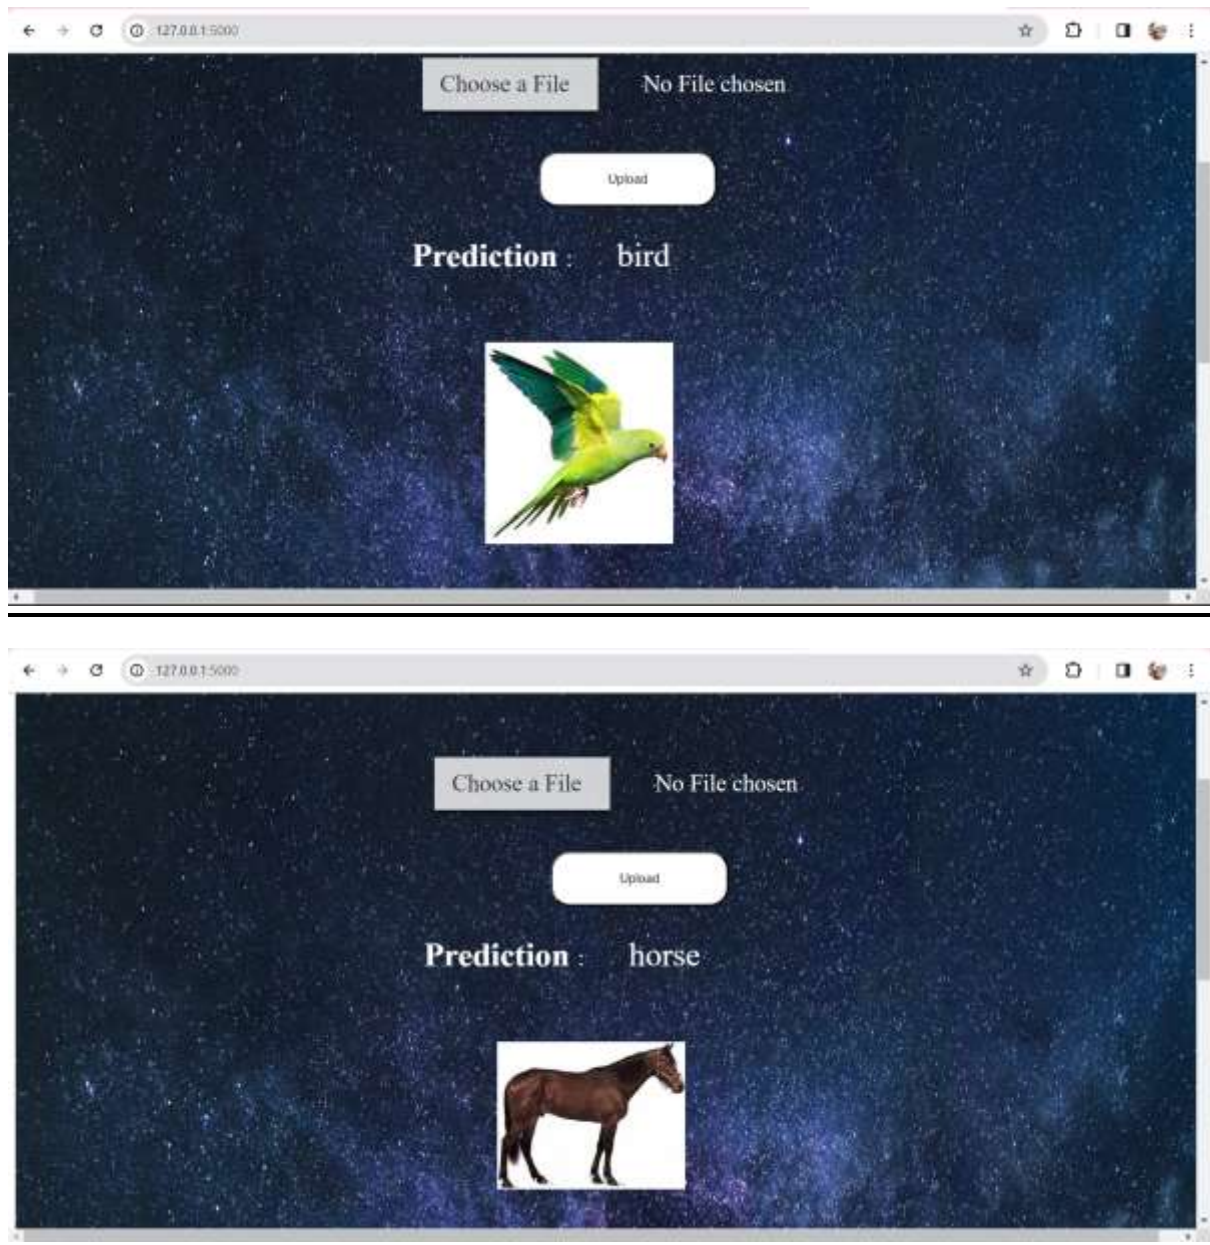
<https://www.figma.com/proto/D3ZmWkJr0Kbd8mInveKXKK/Untitled?type=design&node-id=6-2&t=jPnZa2BCjoHW90ai-0&scaling=scale-down&page-id=0%3A1>

## Screenshots of our Web Page predicting the images



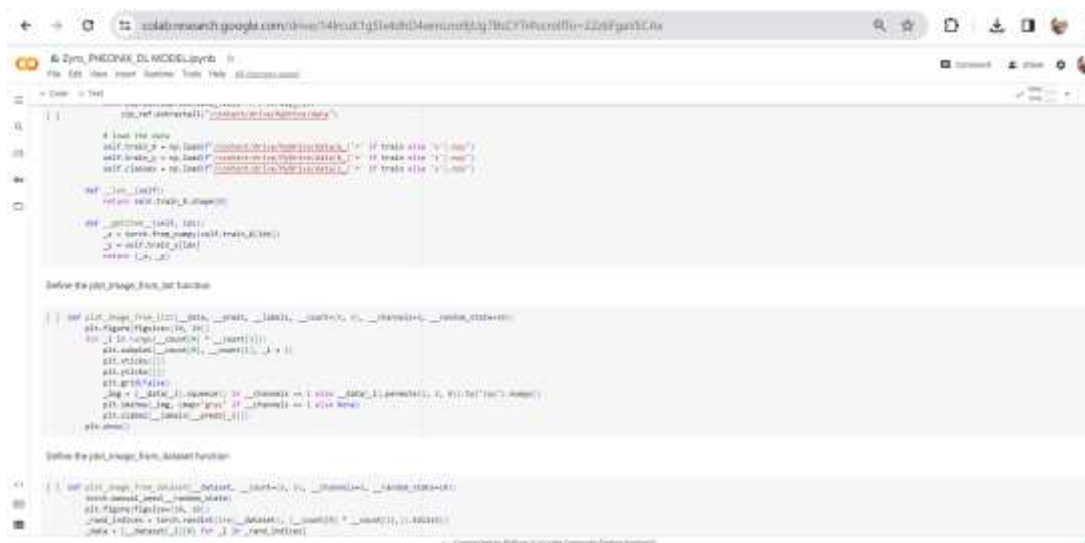
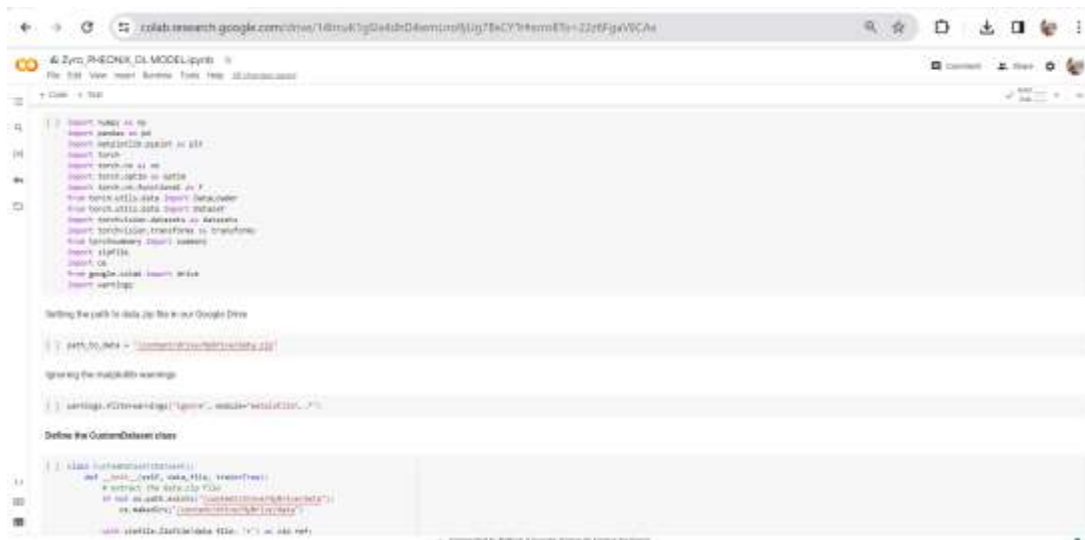
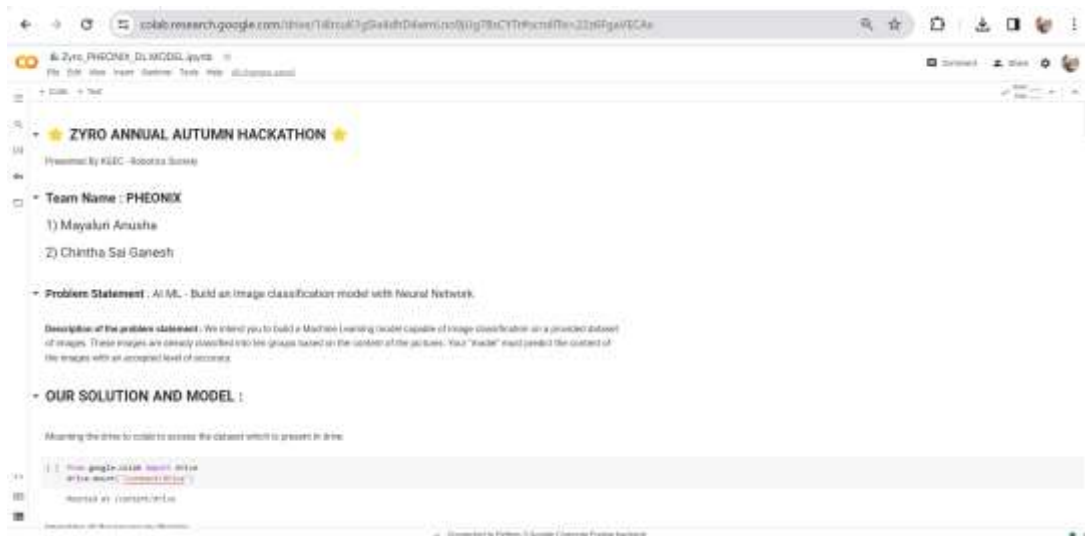


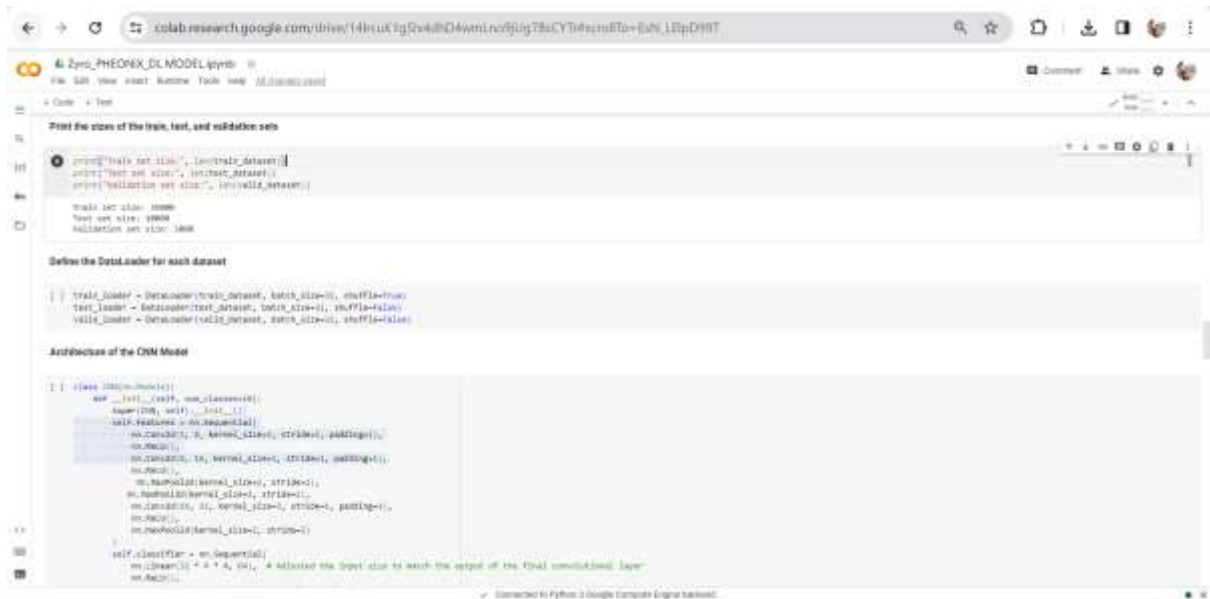
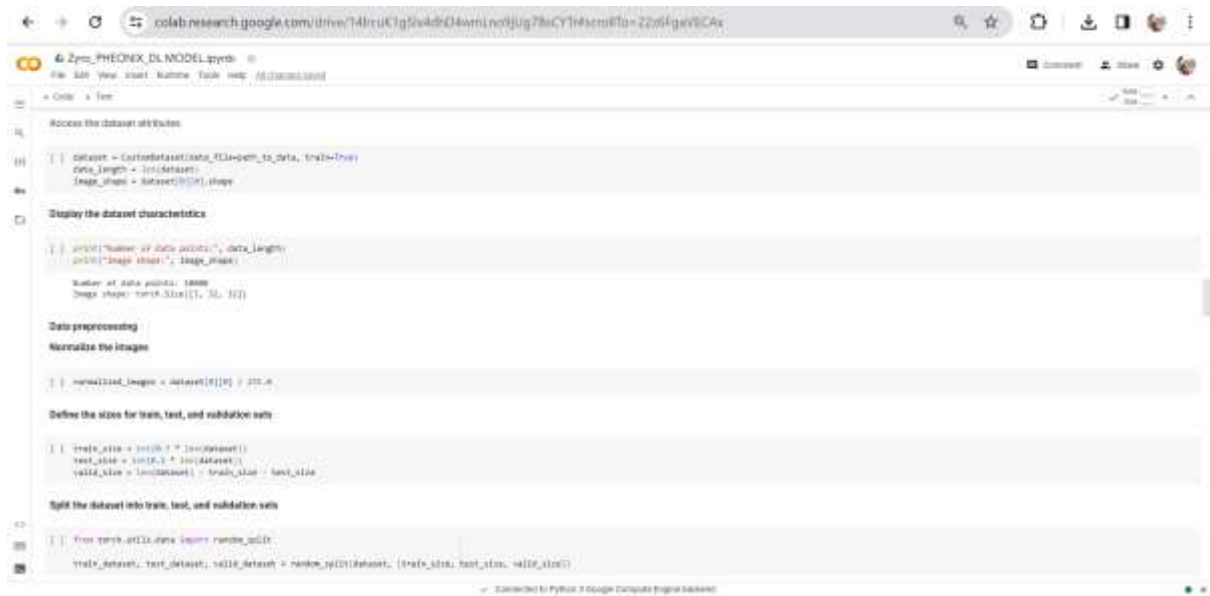





I am delighted to submit my final entry for the Zyro Hackathon's AI ML Image Classifying Challenge. Leveraging PyTorch, I have successfully built a Convolutional Neural Network with a meticulous design that ensures high accuracy, surpassing the minimum requirement of 65 percent, and remarkably keeping the trainable parameters under 2 million. Additionally, I took on the bonus challenge by deploying the model using Flask, creating an interactive web interface for users to effortlessly upload images and receive real-time classifications. The web page is designed for simplicity and functionality, aligning with the competition's criteria of evaluating both accuracy and model lightness. This achievement reflects my commitment to efficiency and innovation in deep learning. The entire project was developed in Google Colab, ensuring accessibility, and the notebook is thoroughly documented with comments for clarity and collaboration.

## Screenshots of the code of the DL Model





```
colab.research.google.com/drive/14lucUg5x4hD4wmiNo9Ug78uCYT9xsm8To=6vX_LBpD997

Zymo_PHEONX_DL_MODEL.pyrb
File Edit View Insert Runtime Tools Help 28/10/2020, 12:00

+ Code + Text

Create an instance of the ModifiedCNN model

11 model = CNN()

Define the loss function and optimizer

12 criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters()) # Adjusted learning rate and added weight decay

Training loop

13 for epoch in range(100): # Increased the number of epochs
    running_loss = 0.0
    for i, data in enumerate(train_loader, 0):
        inputs, labels = data
        optimizer.zero_grad()
        outputs = model(inputs)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()
        running_loss += loss.item()
        if (i + 1) % 100 == 0:
            print('[%d, %5d] Loss: %.3f % 100 epochs = %.2f, running_loss / 100)'
            running_loss = 0.0

    print('Finished Training')

14, 1000 Loss: 0.520
15, 1000 Loss: 0.416
16, 1000 Loss: 0.400
17, 1000 Loss: 0.500
18, 1000 Loss: 0.500
19, 1000 Loss: 0.504
20, 1000 Loss: 0.500
21, 1000 Loss: 0.519
```

```
colab.research.google.com/drive/14lucUg5x4hD4wmiNo9Ug78uCYT9xsm8To=6vX_LBpD997

Zymo_PHEONX_DL_MODEL.pyrb
File Edit View Insert Runtime Tools Help 28/10/2020, 12:00

+ Code + Text

Evaluate the model on the validation set

11 model.eval()
correct = 0
total = 0
with torch.no_grad():
    for data in val_loader:
        inputs, labels = data
        outputs = model(inputs)
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

accuracy = correct / total
print('Validation Accuracy: %.2f%% % accuracy')

Validation accuracy: 70.31%

Evaluate the model on the test set

12 model.eval()
correct = 0
total = 0
with torch.no_grad():
    for data, targets in test_loader:
        inputs = model(inputs)
        _, predicted = torch.max(outputs.data, 1)
        total += targets.size(0)
        correct += (predicted == targets).sum().item()

accuracy = correct / total
print('Test Accuracy: %.2f%%')

Test accuracy: 70.31%

13 model.train()
```

```
colab.research.google.com/drive/14lucUg5x4hD4wmiNo9Ug78uCYT9xsm8To=6vX_LBpD997

Zymo_PHEONX_DL_MODEL.pyrb
File Edit View Insert Runtime Tools Help 28/10/2020, 12:00

+ Code + Text

Print the model summary to get the number of trainable parameters

14 summary(model, input_size=(1, 1, 1))

-----
Layer (type) Output Shape Param #
-----
conv1d-1 [-1, 8, 12, 11] 264
conv1d-2 [-1, 8, 12, 11] 0
conv1d-3 [-1, 16, 12, 11] 5,180
conv1d-4 [-1, 16, 12, 11] 0
conv1d-5 [-1, 16, 16, 16] 0
conv1d-6 [-1, 16, 8, 8] 0
conv1d-7 [-1, 16, 8, 8] 4,040
conv1d-8 [-1, 12, 8, 8] 0
conv1d-9 [-1, 16, 8, 8] 0
conv1d-10 [-1, 16, 8, 8] 10,912
conv1d-11 [-1, 16] 0
conv1d-12 [-1, 16] 0
conv1d-13 [-1, 16] 0
-----
Total params: 16,214
Trainable params: 16,214
Non-trainable params: 0
Input size (MB): 0.01
Forward/backward pass (104 MB): 0.01
Params size (MB): 0.13
Estimated total size (MB): 0.14

Saving the model to integrate flask and create web app

15 torch.save(model.state_dict(), "checkpoint.pth.tar")

Therefore, we have successfully built a model to classify the images into two pre-defined classes with an accuracy of 70.31% and the number of trainable parameters are very less which is equal to 16,214
```



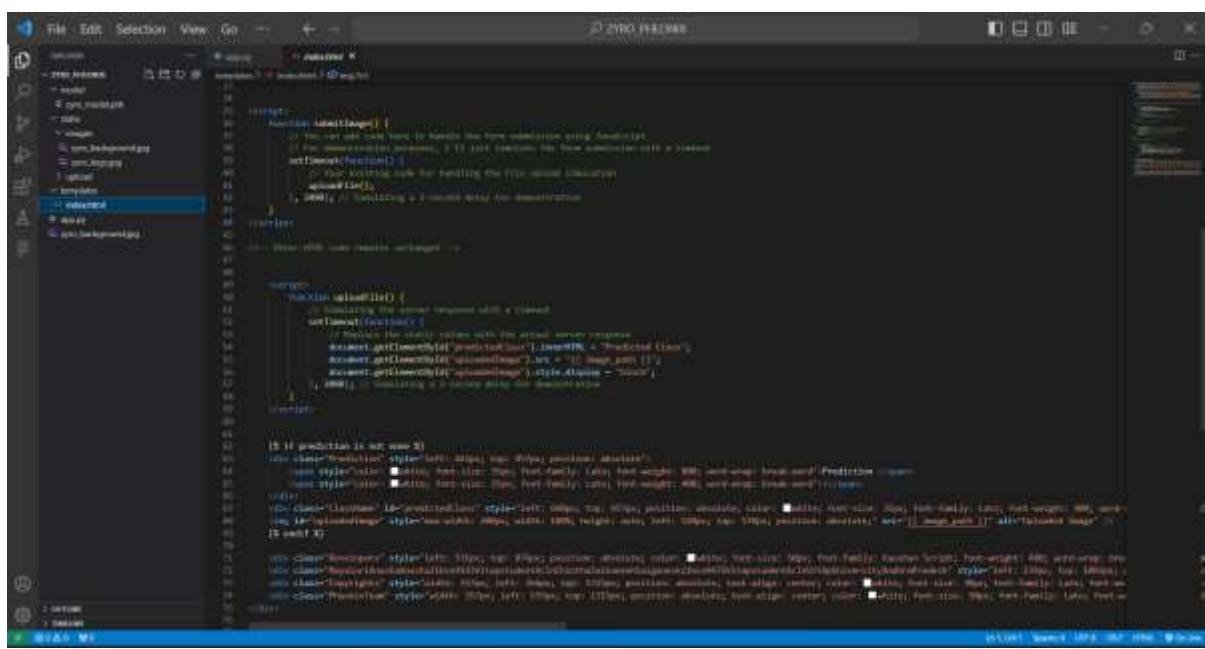
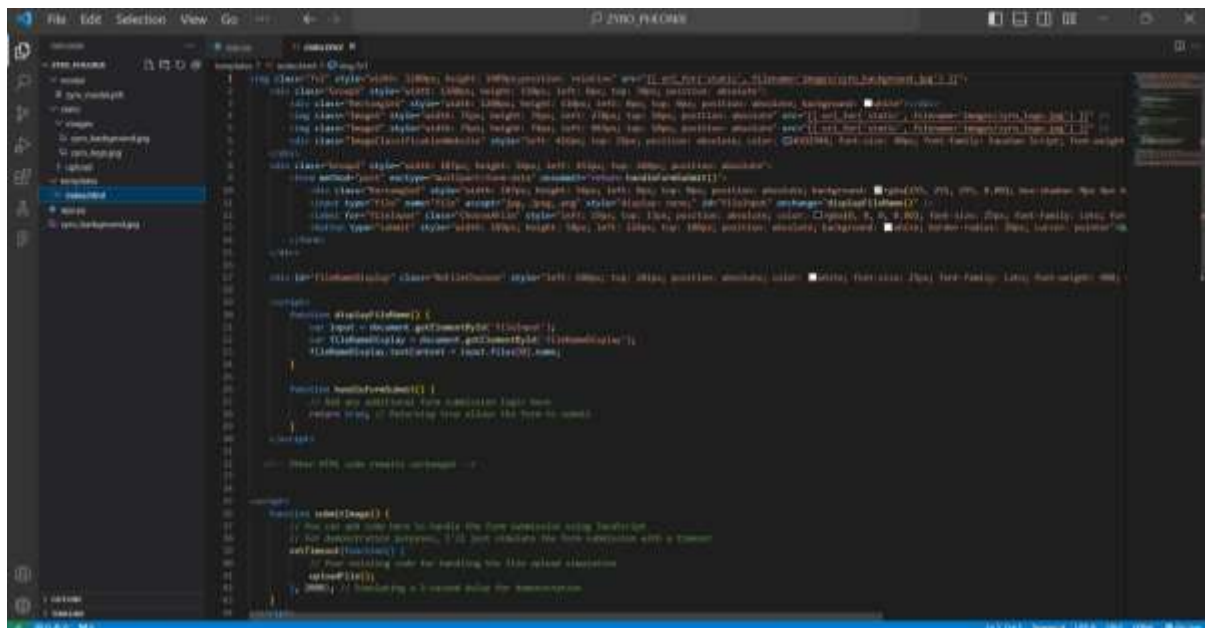


## Screenshots of the code of the Flask App

```
1 from flask import Flask, render_template, request
2 from PIL import Image
3 import numpy as np
4 import torch
5 import torch.nn as nn
6 import os
7 from torchvision import transforms
8
9 app = Flask(__name__)
10 UPLOAD_FOLDER = 'static/upload'
11 app.config['UPLOAD_FOLDER'] = UPLOAD_FOLDER
12
13 import torch.nn as nn
14
15 class CNN(nn.Module):
16     def __init__(self, num_classes=10):
17         super(CNN, self).__init__()
18         self.features = nn.Sequential(
19             nn.Conv2d(3, 8, kernel_size=3, stride=1, padding=1),
20             nn.ReLU(),
21             nn.Conv2d(8, 16, kernel_size=3, stride=1, padding=1),
22             nn.ReLU(),
23             nn.MaxPool2d(kernel_size=2, stride=2),
24             nn.MaxPool2d(kernel_size=2, stride=2),
25             nn.Conv2d(16, 32, kernel_size=3, stride=1, padding=1),
26             nn.ReLU(),
27             nn.MaxPool2d(kernel_size=2, stride=2)
28         )
29         self.classifier = nn.Sequential(
30             nn.Linear(32 * 4 * 4, 64), # Adjusted the input size to match the output of the final convolutional layer
31             nn.ReLU(),
32             nn.Dropout(0.5), # Added dropout for regularization
33             nn.Linear(64, num_classes)
34         )
35
36     def forward(self, x):
37         x = self.features(x)
38         x = x.view(x.size(0), -1)
39         x = self.classifier(x)
40         return x
```

```
38         x = self.features(x)
39         x = x.view(x.size(0), -1)
40         x = self.classifier(x)
41         return x
42
43 # Create an instance of the CNN model
44 model = CNN()
45 model.load_state_dict(torch.load('model/cyru_model.pth'))
46 model.eval()
47
48 # Define image preprocessing
49 preprocess = transforms.Compose([
50     transforms.Resize((32, 32)),
51     transforms.ToTensor(),
52 ])
53
54 def predict_image(image_path):
55     img = Image.open(image_path).convert('RGB')
56     img = preprocess(img)
57     img = img.unsqueeze(0) # Add batch dimension
58     with torch.no_grad():
59         output = model(img)
60     predicted_class = torch.max(output, 1)
61     class_names = {0: 'airplane', 1: 'automobile', 2: 'bird', 3: 'cat', 4: 'deer',
62                   5: 'dog', 6: 'frog', 7: 'horse', 8: 'ship', 9: 'truck'}
63     class_name = class_names[predicted_class.item()]
64     return class_name
65
66 @app.route('/', methods=['GET', 'POST'])
67 def index():
68     if request.method == 'POST':
69         file = request.files['file']
70         if file:
71             file_path = os.path.join(app.config['UPLOAD_FOLDER'], file.filename)
72             file.save(file_path)
73             prediction = predict_image(file_path)
74             return render_template('index.html', prediction=prediction, image_path=file_path)
75
76     return render_template('index.html', prediction=None, image_path=None)
77
78 if __name__ == '__main__':
79     app.run(debug=True)
```

### Screenshots of the code of the UI of the Web Page



## Workflow Video :

[https://drive.google.com/file/d/1e-TxUcSdsepUxiPpN3lkfrJr0DKxft\\_/view?usp=sharing](https://drive.google.com/file/d/1e-TxUcSdsepUxiPpN3lkfrJr0DKxft_/view?usp=sharing)

## Highlights of our project is are:

- Very smaller number of parameters . i.e., 39514

```
=====
Total params: 39,514
Trainable params: 39,514
Non-trainable params: 0
-----
Input size (MB): 0.01
Forward/backward pass size (MB): 0.45
Params size (MB): 0.15
Estimated Total Size (MB): 0.61
-----
```

- Industry acceptable range accuracy. i.e., 70.39%

### Evaluate the model on the test set

```
[ ] model.eval()
    with torch.no_grad():
        correct = 0
        total = 0
        for data, targets in test_loader:
            outputs = model(data)
            _, predicted = torch.max(outputs.data, 1)
            total += targets.size(0)
            correct += (predicted == targets).sum().item()

    accuracy = 100 * correct / total
    print(f"Test Accuracy: {accuracy:.2f}%")
```

Test Accuracy: 70.39%



➤ Very User-Friendly Web Page



➤ Summary of the model

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 8, 32, 32]	224
ReLU-2	[-1, 8, 32, 32]	0
Conv2d-3	[-1, 16, 32, 32]	1,168
ReLU-4	[-1, 16, 32, 32]	0
MaxPool2d-5	[-1, 16, 16, 16]	0
MaxPool2d-6	[-1, 16, 8, 8]	0
Conv2d-7	[-1, 32, 8, 8]	4,640
ReLU-8	[-1, 32, 8, 8]	0
MaxPool2d-9	[-1, 32, 4, 4]	0
Linear-10	[-1, 64]	32,832
ReLU-11	[-1, 64]	0
Dropout-12	[-1, 64]	0
Linear-13	[-1, 10]	650

## ➤ Model Architecture

```
CNN(  
  (features): Sequential(  
    (0): Conv2d(3, 8, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (1): ReLU()  
    (2): Conv2d(8, 16, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (3): ReLU()  
    (4): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (5): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (6): Conv2d(16, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (7): ReLU()  
    (8): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
  )  
  (classifier): Sequential(  
    (0): Linear(in_features=512, out_features=64, bias=True)  
    (1): ReLU()  
    (2): Dropout(p=0.5, inplace=False)  
    (3): Linear(in_features=64, out_features=10, bias=True)  
  )  
)
```

**\*\*\*\*THANK YOU\*\*\*\***